


Auditing the auditor: the same answers scored by each verifier



Source: eval/results/verifier_comparison.json — * p<0.05 ** p<0.01 *** p<0.001 (permutation tests, 20k resamples); specificity n = questions with a grounded operand sentence