


The labeled failure-case dataset (Contribution 2 deliverable)



- false negative (correct + grounded, scored unfaithful)
- false positive (wrong / ungrounded, scored faithful)
- insensitive (targeted removal: no drop)
- nonspecific (targeted drop $\leq$ random drop)
- unscorable (no statements)